

Abstraction: Process (cont.)





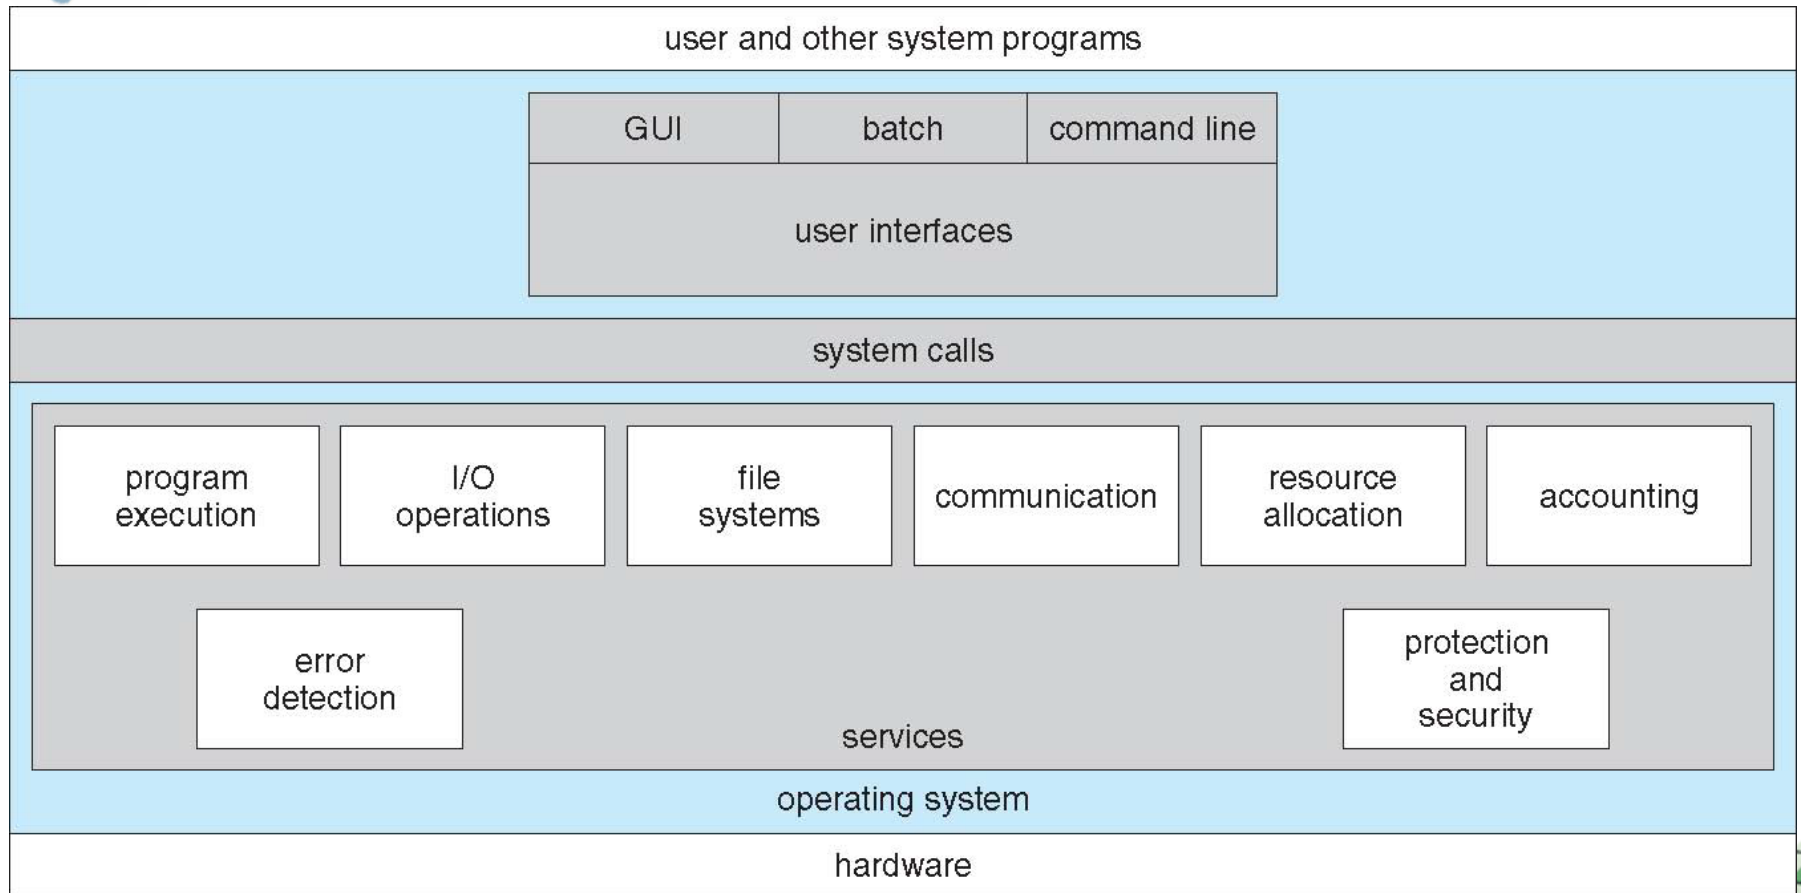
Objectives


- How are processes represented and controlled by the OS.
- **Process states** which characterize the behaviour of processes.
- **Data structures** used to manage processes.
- Ways in which the OS uses these data structures to control process execution.





Operating System Services






Types of System Calls

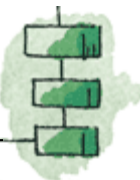
- Process control
 - load, execute; create process, terminate process; get process attributes, set process attributes; wait for time, wait for events; allocate and free memory; end, abort;
- File management
 - ...
- Device management
 - ...
- Information maintenance
 - ...
- Communications
 - ...






Dual Mode Operation

- If every program had unfettered access to system resources, a program could look at all memory locations, including that of other programs, as well as read all the data on all of the attached disks
- Dual-mode operation is then designed to provide a layer of protection and stability to computer systems by separating user programs and the operating system into two modes: user mode and kernel mode
- User mode restricts access to privileged resources, while kernel mode has full access to these resources





Dual Mode Operation


User Mode

- user program executes in user mode
- certain areas of memory are protected from user access
- certain instructions may not be executed

Kernel Mode

- OS executes in kernel mode
- privileged instructions may be executed
- protected areas of memory may be accessed





Dual Mode Operation


- The advantages of dual mode operation include
 - provides a layer of protection between user programs and OS to protect the system from unauthorized and malicious access
 - Reduce the risk of malware attacks or other security threats as user programs in user mode cannot modify system data or perform privileged operations
 - Ensure system-level operations are performed correctly and reliably
 - Prevent user programs from accidentally or maliciously causing system crashes or other errors






System Calls

- Typically, a number associated with each system call
 - System-call interface maintains a table indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call
 - Most details of OS interface hidden from programmer by API
 - Managed by run-time support library (set of functions built into libraries included with compiler)



System Calls


- It should be obvious that we don't want user-mode programs to easily switch to kernel mode
 - that would make the two privilege levels useless
- However, a user-mode program does need to cross into kernel mode
 - E.g., when a program wants to read from disk
- We need a mechanism whereby
 - a user-mode program can switch into kernel mode
 - but have no control over the instructions which will be performed in kernel mode





System Calls

- The user program
 - sets up the parameters
 - traps to the kernel
 - Different architectures have different trap instructions, e.g., int on x86
- System call enters the kernel at a fixed location
 - Switches the stack pointer SP to the kernel stack
 - Saves the user SP value
 - Saves the user PC value (= return address)
 - Saves the user privilege mode
 - Sets the kernel privilege mode
 - Sets the new PC to the system call handler





System Calls

- Kernel system call handler carries out the desired system call
 - Saves user registers
 - Examines the system call number
 - Checks arguments
 - Performs operation
 - Stores result
 - Performs a “return from system call” instruction
 - Restores user privilege mode, SP and PC and registers if returning directly to the user process which made the call
 - Or go through the dispatcher to select the next process to run
- All memory used by the system call handler is in kernel space, so it is protected from interference by the user program



System Calls

Process P



RAM

P wants to call read()



System Calls


Process P



RAM

P can only see its own memory because of **user mode**
(other areas, including kernel, are hidden)

P wants to call `read()` but no way to call it directly



System Calls

Process P




RAM

```
movl $6, %eax;    int $64
```





System Calls




```
movl $6, %eax;    int $64
```



System Calls



System Calls



```
movl $6, %eax;    int $64
```

Follow entries to correct system call code



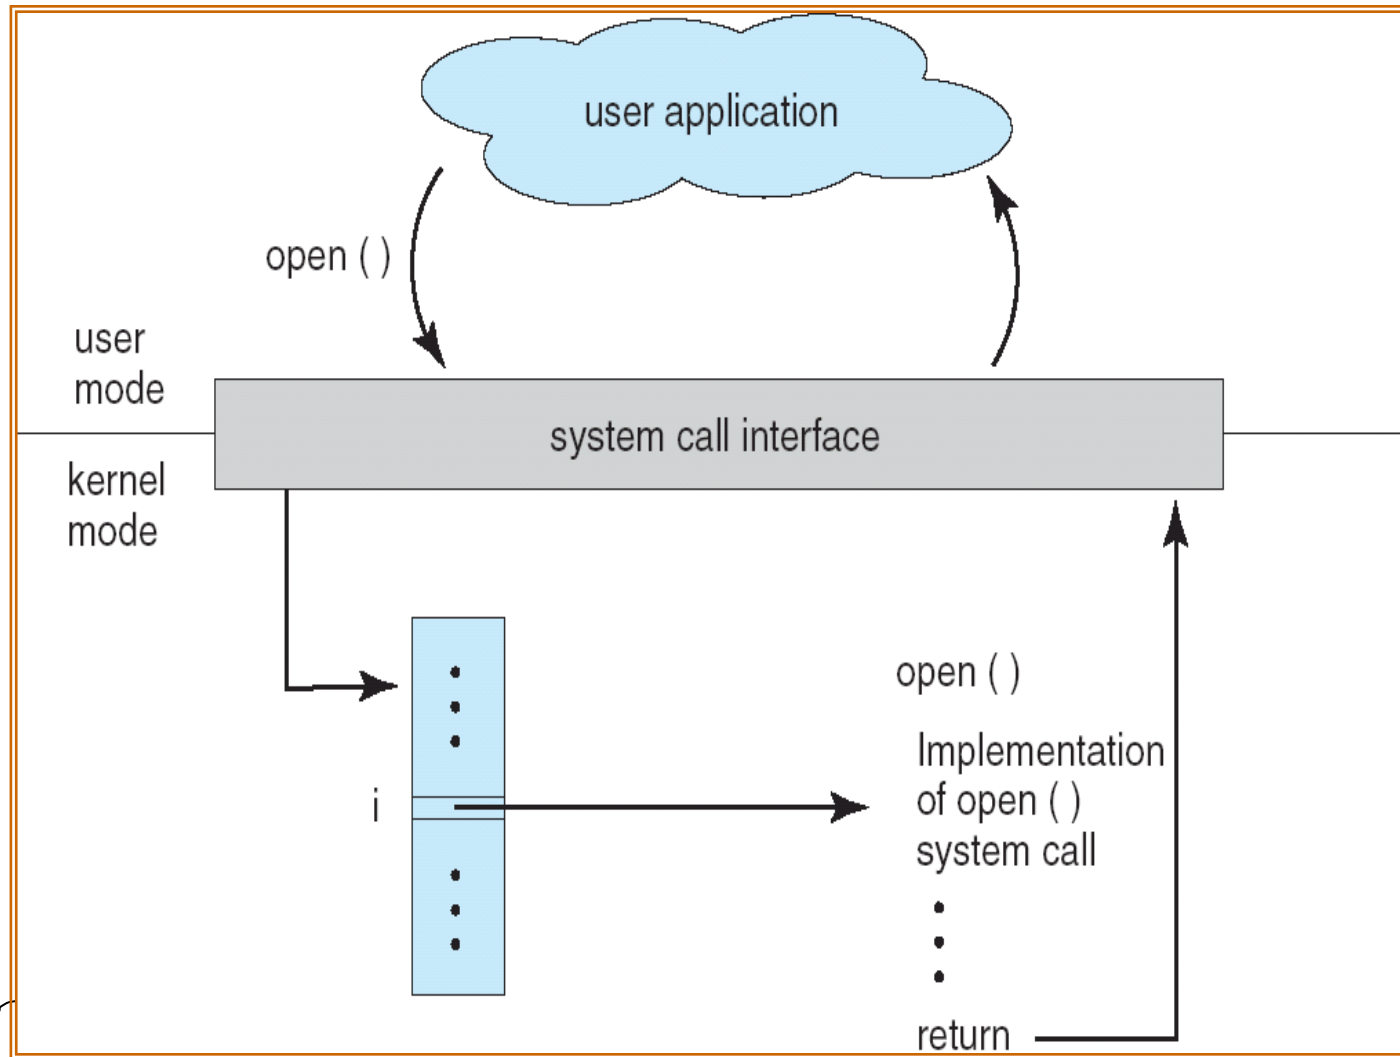



System Calls

- System calls are mostly accessed by user programs via a high-level Application Programming Interface (API) rather than direct system call use
 - E.g., C library implements standard API from kernel system call interface



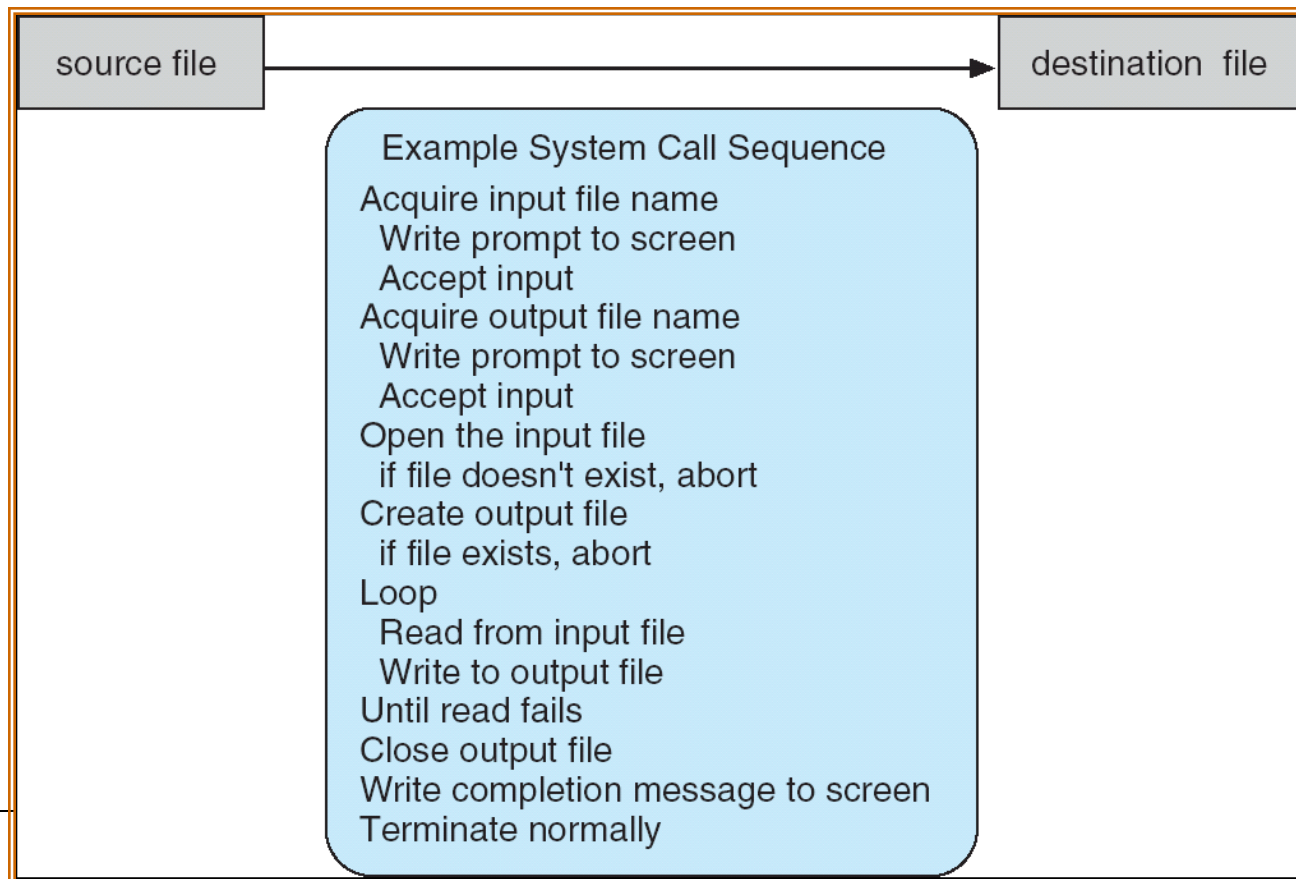
API – System Call – OS Relationship






Example of System Calls

- System call sequence to copy the contents of one file to another file





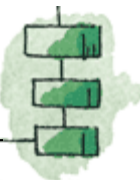

Processes and Threads

- Processes have two characteristics:
 - **Resource ownership** - process includes a virtual address space to hold the process image
 - the OS performs a protection function to prevent unwanted interference between processes with respect to resources
 - **Scheduling/execution** - follows an execution path that may be interleaved with other processes
 - a process has an execution state (Running, Ready, etc.) and a dispatching priority and is scheduled and dispatched by the OS
- These two characteristics are treated independently by modern operating systems:
 - The unit of dispatching is referred to as a **thread** or lightweight process
 - The unit of resource ownership is referred to as a **process** or **task**



Process

- The unit of resource allocation and a unit of protection
- A process is associated with
 - A virtual address space which holds the process image
 - Protected access to
 - Processors,
 - Other processes,
 - Files,
 - I/O resources





Multiple Threads in Process


- Each thread has
 - Access to the memory and resources of its process (all threads of a process share this)
 - An execution state (running, ready, etc.)
 - Saved thread context when not running
 - An execution stack
 - Some per-thread static storage for local variables







Threads vs. processes

Single-Threaded Process Model



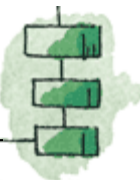
Multithreaded Process Model






Thread Synchronization


- It is necessary to synchronize the activities of the various threads in a process so that they do not interfere with each other
 - all threads of a process share the same address space and other resources
 - any alteration of a resource by one thread affects the other threads in the same process
 - Synchronization needed to coordinate execution of multiple threads



Linux Threads

- Linux uses the same internal representation for processes and threads; a thread is simply a new process (or task) that happens to share the same address space as its parent
- A distinction is only made when a new thread is created by the `clone` system call
 - `fork` creates a new process with its own entirely new process context
 - `clone` creates a new process with its own identity, but that is allowed to share the data structures of its parent






Linux Threads

- Using `clone` gives an application fine-grained control over exactly what is shared between two threads

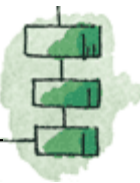
flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.





Windows Threads

- Windows makes use of two types of process-related objects:
- Processes
 - an entity corresponding to a user job or application that owns resources
- Threads
 - a dispatchable unit of work that executes sequentially and is interruptible



Summary

- Operating System services
 - System calls
 - Dual mode to separate user-mode from kernel mode for security
 - Syscall: call kernel mode functions
 - Transfer from user-mode to kernel-mode (trap)
 - Return from kernel-mode to user-mode (return-from-trap)
- Threads
 - Process related to resource ownership
 - Thread related to program execution
 - Threads of a process share the same address space and resources
 - Any alteration of a resource by one thread affects the other threads in the same process

