





Concurrency: Mutual Exclusion and Synchronization

Outline

- Race condition
- Critical section
- Mutual exclusion
- Hardware support
 - Atomic operations
 - Special machine instructions
 - Compare&Swap
 - Exchange





Multiple Processes/Threads

- The design of modern Operating Systems is concerned with the management of multiple processes and threads
 - Multiprogramming
 - Multiprocessing
- Big Issue is Concurrency
 - Managing the interaction of processes/threads






Race Condition

- A race condition (or data race) occurs when
 - Multiple processes or threads read and write shared data items
 - They do so in a way where the final result depends on the order of execution of the processes/threads
 - The output depends on who finishes the race last






A Simple Example on a Single Processor System

- `count++` could be implemented as
 `register1 = count`
 `register1 = register1 + 1`
 `count = register1`
- `count--` could be implemented as
 `register2 = count`
 `register2 = register2 - 1`
 `count = register2`





A Simple Example on a Single Processor System

- Consider:
 - **process A** increments count and **process B** decrements count simultaneously
 - the execution interleaving with “count = 5” initially:

S0: **process A** execute **register1 = count** {register1 = 5}
S1: **process A** execute **register1 = register1 + 1** {register1 = 6}
S2: **process B** execute **register2 = count** {register2 = 5}
S3: **process B** execute **register2 = register2 - 1** {register2 = 4}
S4: **process A** execute **count = register1** {count = 6}
S5: **process B** execute **count = register2** {count = 4}





Critical Section

- In the count example, the correct result will be guaranteed if one process starts increment/decrement only after the other one has finished decrement/increment
 - Need to have mechanism, i.e. mutual exclusion - lock/unlock to serialize the processes/threads entering their critical sections
- When a process executes code that manipulates shared data (or resource), we say that the process is in its Critical Section
 - Want to make sure only one process in its critical section at any time
- A general structure:

...


entry section

critical section

exit section

noncritical section

...






Atomic Operation

- A sequence of instructions appears to be indivisible
- The idea behind making a series of actions atomic is simply expressed with the phrase “all or nothing”
- It should either appear as if
 - all of the actions you wish to group together occurred
 - or that none of them occurred
 - with no in-between state visible
- Sometimes, the grouping of many actions into a single atomic action is called a transaction
 - an idea developed in great detail in the world of databases and transaction processing






Mutual Exclusion

- Only one process at a time is allowed in the critical section for a resource
- No assumptions are made about relative process speeds or number of processes
- For mutual exclusion to work we need
 - hardware support – to provide basic atomic operation primitives
 - OS – to ensure efficiency



Mutual Exclusion: Hardware Support

- Interrupt Disabling
 - A process runs until it invokes an operating system service or until it is interrupted
 - Disabling interrupts guarantees mutual exclusion
 - Work in uniprocessor systems
- Disadvantages:
 - The efficiency of execution could be noticeably degraded
 - This approach will not work in a multiprocessor architecture
 - Serious security issues – thus user process/threads must never disable interrupts





Mutual Exclusion: Hardware Support

- Special Machine Instructions:
 - Compare&Swap Instruction
 - also called a “compare and exchange instruction”
 - Exchange Instruction
- These are atomic instructions
 - Operations are indivisible





Compare&Swap Instruction

```
int compare_and_swap (int *word,  
    int testval, int newval)  
{  
    =0                =1  
    int oldval;  
    oldval = *word;  
    if (oldval == testval) *word = newval;  
    return oldval;  
}
```

- If word = 1, unchange, and return 1
- If word = 0, word = 1, and return 0








Compare&Swap Instruction

```
/* program mutual exclusion */
const int n = /* number of processes */;
int bolt;
void P(int i)
{
    while (true) {
        while (compare_and_swap(bolt, 0, 1) == 1)
            /* do nothing */;
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main()
{
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

Busy waiting



(a) Compare and swap instruction



Exchange instruction

```
void exchange (int register, int  
memory)  
{  
    int temp;  
    temp = memory;  
    memory = register;  
    register = temp;  
}
```





Exchange Instruction

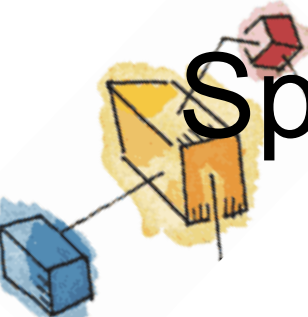
```
/* program mutualexclusion */
int const n = /* number of processes**/;
int bolt;
void P(int i)
{
    int keyi = 1;
    while (true) {
        do exchange (keyi, bolt)
        while (keyi != 0);
        /* critical section */;
        bolt = 0;
        /* remainder */;
    }
}
void main()
{
    bolt = 0;
    parbegin (P(1), P(2), ..., P(n));
}
```

Busy waiting ←



(b) Exchange instruction






Special Machine Instructions: Advantages

- Applicable to any number of processes on either a single processor or multiple processors sharing main memory
- It is simple and therefore easy to verify
- It can be used to support multiple critical sections; each critical section can be defined by its own variable





Special Machine Instructions:

Disadvantages

- Busy-waiting is employed, thus while a process is waiting for access to a critical section it continues to consume processor time
- Starvation is possible when a process leaves a critical section and more than one process is waiting
 - Some process could indefinitely be denied access.
- Deadlock is possible

