




Processor Scheduling



Outline

- Types of scheduling
 - Long term, medium term and short term
- Scheduling criteria
 - Turnaround time
 - Response time
- Scheduling policies/algorithms
 - Preemptive and non-preemptive
 - FCFS
 - SJF
 - SRT
 - RR
 - MFBQ






Processor Scheduling

- The resource provided by a processor is execution time
 - The resource is allocated by means of a schedule
- Aim is to assign processes to be executed by the processor in a way that meets system objectives, such as turnround time, response time, and processor efficiency






Types of Scheduling

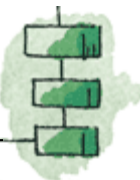
- Broken down into three separate functions:






Long-Term Scheduling

- Determines which programs are admitted to the system for processing
- Controls the degree of multiprogramming
 - More processes, smaller percentage of time each process is executed
 - may limit to provide satisfactory service to the current set of processes






Medium-Term Scheduling


- Part of the swapping function
- Swapping-in decisions are based on the need to manage the degree of multiprogramming
- considers the memory requirements of the swapped-out processes






Short-Term Scheduling

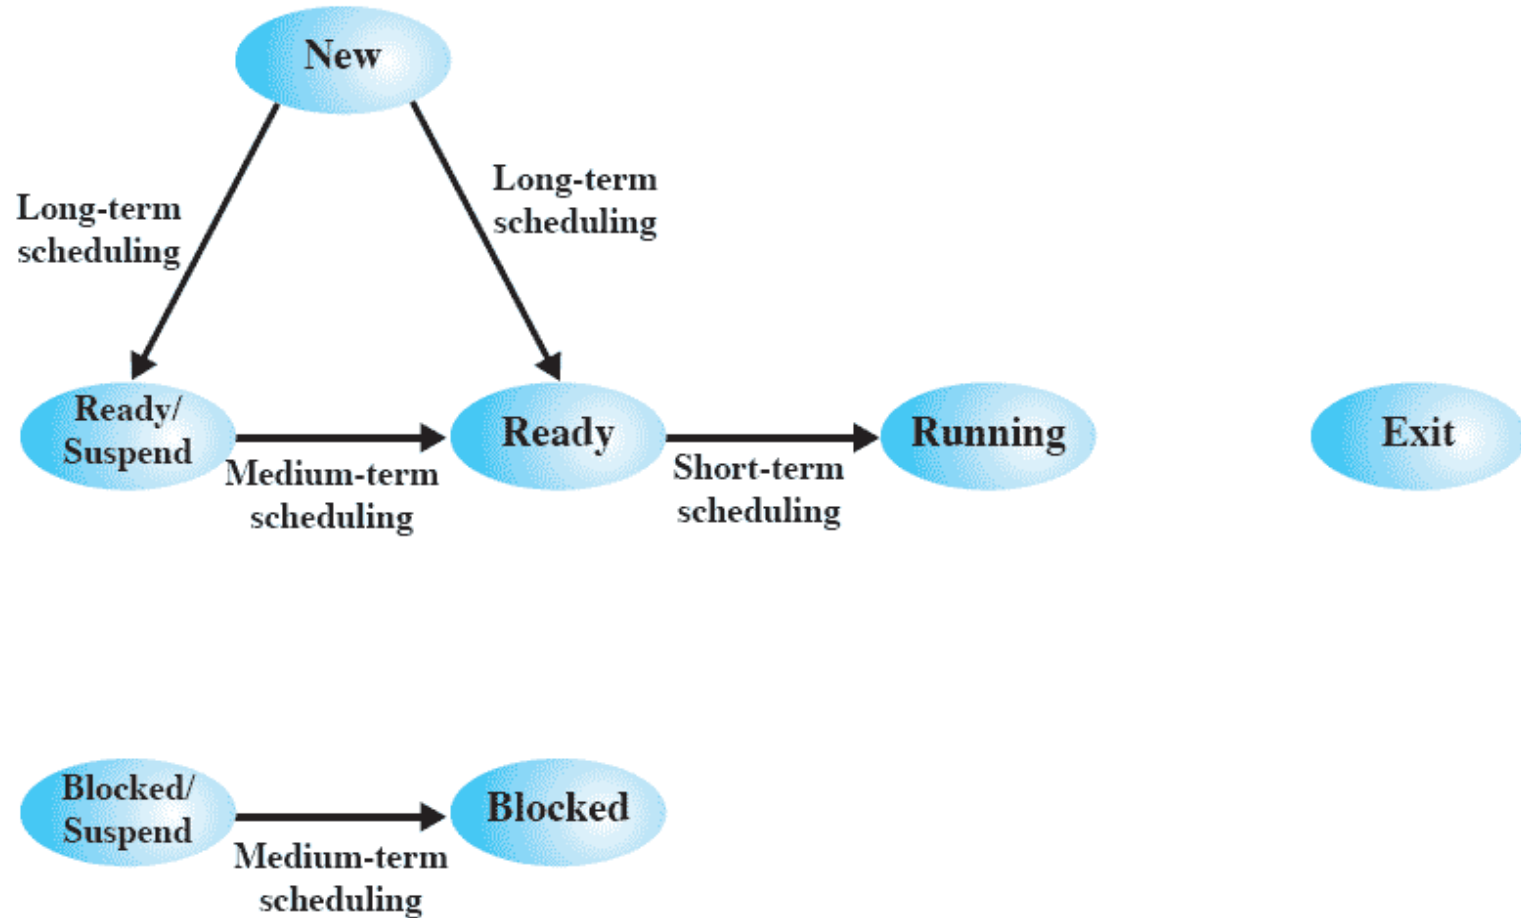
- Known as the dispatcher
- Executes most frequently
- Invoked when an event occurs
 - Clock interrupts
 - I/O interrupts
 - Operating system calls
 - Signals



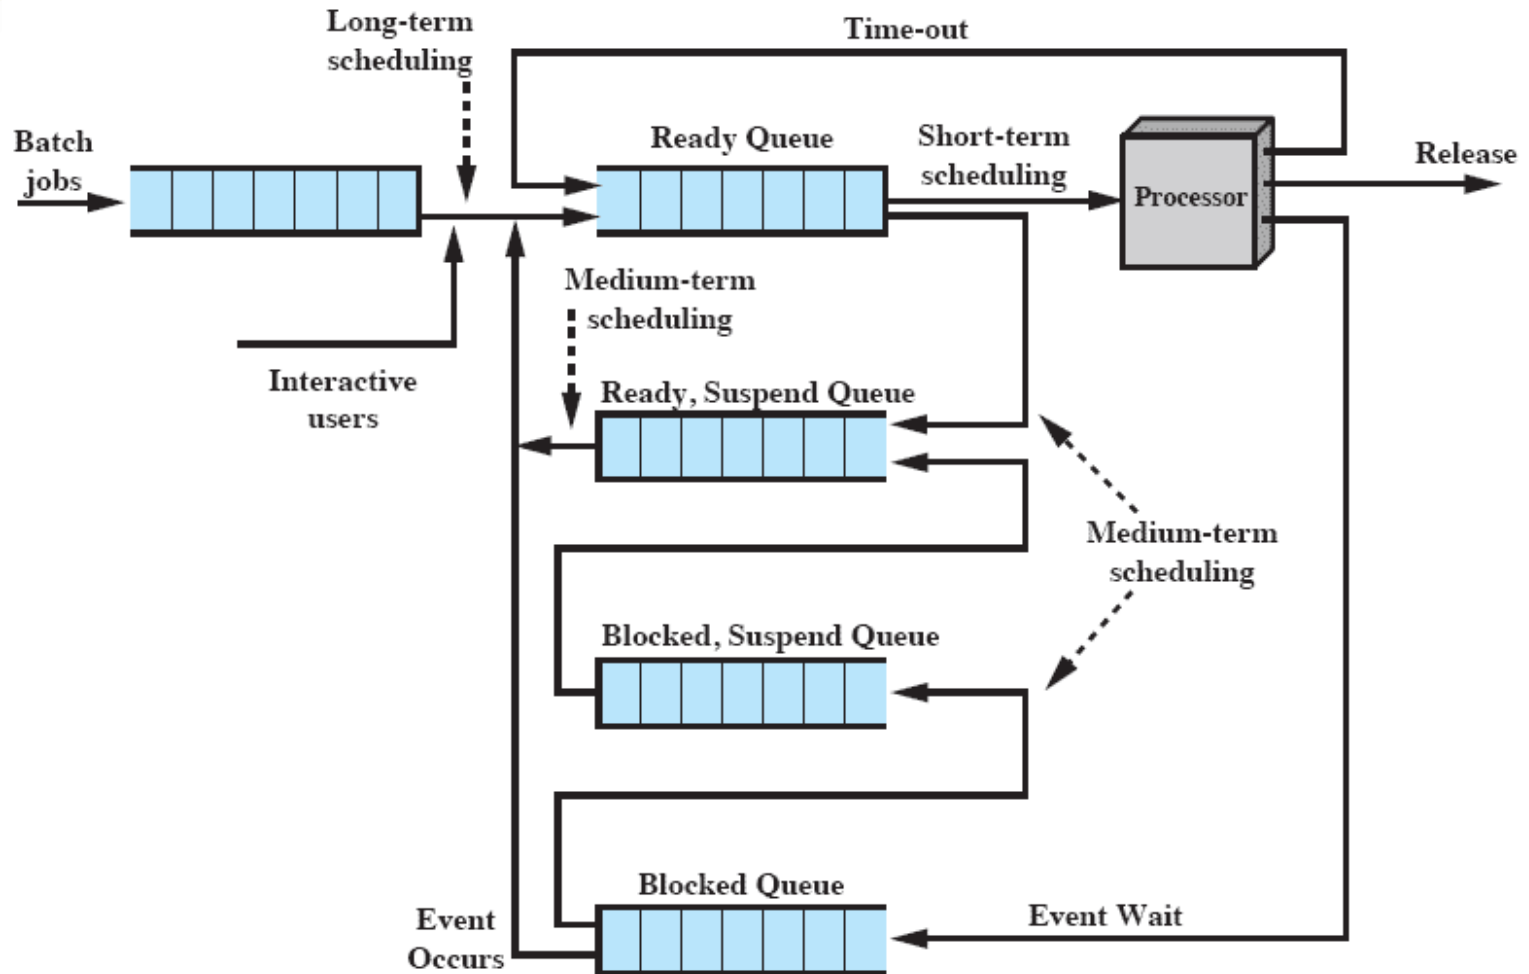
Two Suspend States



Scheduling and Process State Transitions



Queuing Diagram






Aim of Short Term Scheduling

- Main objective is to allocate processor time to optimize certain aspects of system behaviour
- A set of criteria is needed to evaluate the scheduling policy
 - e.g., turnaround time, response time, deadlines, fairness






Process Scheduling Example

- Assume a set of processes

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

- Service time is the required total execution time
- We want to calculate turnaround time & response time






Turnaround Time

- $\text{Turnaround_time} = \text{completion_time} - \text{arrival_time}$
 - Including both execution_time and wait_time
 - $\text{Wait_time} = \text{turnaround_time} - \text{execution_time}$
- Example:
 - A process's execution time is 15
 - It arrives at time 10 and completed execution at time 30
 - Then the $\text{turnaround_time} = 30 - 10 = 20$
 - $\text{Wait_time} = 20 - 15 = 5$






First-Come-First-Served

- Each process joins the Ready queue in a first-in-first-out manner
- When the current process ceases to execute, the process waiting the longest time in the Ready queue is selected


Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

First-Come-First Served (FCFS)



Turnaround time = Finish time - Arrival time
Wait time = Turnaround time - Service time






First-Come-First-Served

- Process A:
 - Turnaround_time = $3 - 0 = 3$
 - Wait_time = $3 - 3 = 0$
- Process B:
 - Turnaround_time = $9 - 2 = 7$
 - Wait_time = $7 - 6 = 1$
- Process C:
 - Turnaround_time = $13 - 4 = 9$
 - Wait_time = $9 - 4 = 5$
- Process D:
 - Turnaround_time = $18 - 6 = 12$
 - Wait_time = $12 - 5 = 7$
- Process E:
 - Turnaround_time = $20 - 8 = 12$
 - Wait_time = $12 - 2 = 10$

Average turnaround_time =
 $(3+7+9+12+12)/5 = 8.6$


Average wait_time =
 $(0+1+5+7+10)/5 = 4.6$






First-Come-First-Served

- The simplest scheduling policy
- Issues:
 - A short process may have to wait a very long time before it can execute
 - Favors CPU-bound processes
 - I/O-bound processes have to wait until CPU-bound process completes







Shortest Job First (SJF)

- or Shortest Process Next (SPN)
- Process with shortest expected processing time is selected next
- Short process jumps ahead of longer processes

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2


**Shortest Job
First (SJF)**






Shortest Job First (SJF)

- Average turnaround_time = 7.6
- Average wait_time = 3.6
- Favors short jobs
 - shorter turnaround and wait time
- Issues:
 - Predictability of longer processes is reduced
 - Possibility of starvation for longer processes






Non-preemptive vs Preemptive

- Non-preemptive
 - Once a process is in the running state, it will continue until it terminates or blocks itself for I/O
 - FCFS and SJF are non-preemptive policies
 - Even using SJF short jobs still need to wait for a long time if a very long job has been scheduled and running
- Preemptive
 - Currently running process may be interrupted and moved to ready state by the OS



- 





Shortest Remaining Time (SRT)


- Average turnaround_time = ?
- Average wait_time = ?
- Should give superior turnaround time performance to SJF because a short job is given immediate preference to a running longer job
- However, risk of starvation of longer processes
- Both SJF and SRT policies
 - must estimate processing time and choose the shortest
 - Question: how?
 - Possibility of starvation for longer processes






Response Time

- $\text{Response_time} = \text{first_run_time} - \text{arrival_time}$



- B's turnaround_time = 20
- B's response_time = 10






Round Robin (RR)

- Clock interrupt is generated at periodic intervals
- When an interrupt occurs, the currently running process is placed in the ready queue
 - Next ready job is selected






Round Robin (RR)

- Uses preemption based on a clock
 - also known as time slicing, because each process is given a slice of time before being preempted.


Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

Round-Robin
(RR), $q = 1$



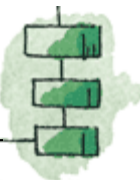
Wait time = ?








Round Robin (RR)

- Round robin scheduling policy
 - Decreases response time
 - But will significantly increase turnaround time when multiple long jobs are running simultaneously
- Tuning challenges:
 - What is a good time slice for round robin?
 - What is the overhead of context switching?




Round Robin (RR)





Average Response Time?

Average Turnaround Time?





Round Robin (RR)



Average Response Time

$$(0 + 5 + 10)/3 = 5s$$

$$(0 + 1 + 2)/3 = 1s$$

Average Turnaround Time

$$(5 + 10 + 15)/3 = 10s$$

$$(13 + 14 + 15)/3 = 14s$$





Next Week's Quiz

- **Next week's quiz:**
- Assume three processes A, B and C
 - Arrived at the same time (say at time 0)
 - A's execution time is 20
 - B's execution time is 1.2
 - C's execution time is 1
- A starts first, then B and C, calculating each process's turnaround time and wait time, and then average turnaround time and average wait time When FCFS is used
- When RR is used (with time slice = 1), calculate turnaround times, wait times and averages
- Compare and discuss the results
- What conclusions can you make from the results?






Multilevel Feed Back Queue (MFBBQ)


- To optimize turnaround time
 - Running shorter jobs first
 - Need a priori knowledge of each job's length – How?
- To minimize response time
 - Round robin
 - Sometimes terrible for turnaround time, especially when multiple long jobs run simultaneously
- How can we design a scheduler that minimizes response time for interactive jobs while also minimizing turnaround time without a priori knowledge of job length?



Multilevel Feed Back Queue (MFBQ)


- A number of distinct queues
- Each assigned a different priority level
- At any given time a job that is ready to run is on a single queue





Multilevel Feed Back Queue (MFBBQ)


- How to select a job to run?
 - Always choose a job with higher priority to run
 - **If $\text{priority}(A) > \text{Priority}(B)$, A runs**
 - **If $\text{priority}(A) == \text{Priority}(B)$, A & B run in RR**
- How to set priority to each job?
 - Dynamically vary the priority of a job based on its behaviour
 - Use the history of the job to predict its future behaviour
 - **Jobs start at top priority**
 - Each job is given a time slice to run
 - **If job uses the whole slice, demote process to the next lower priority queue**





Multilevel Feed Back Queue (MFBBQ)

- Issue:
 - Since the CPU is always allocated to the process with the highest priority, low priority processes at the bottom queue may never execute – **starvation!**
- Solution:
 - After some time period S , move all the jobs in the system to the topmost queue
 - Promoted job will share the CPU with other high-priority jobs in a round-robin fashion
 - Guaranteed not to starve






Multilevel Feed Back Queue (MFBQ)

- How to parameterize MFBQ?
 - How many queues should there be?
 - How big should the time slice be per queue?
 - How often should priority be boosted in order to avoid starvation?
- No easy answers
 - Workloads dependent





Contemporary Scheduling

- CPU sharing -- timer interrupts
 - Time quantum (or time slice) determined by interval timer
- With preemption
- Priority-based process (job) selection
 - Select the highest priority process
 - Priority reflects policy
- Usually a variant of Multilevel Feedback Queues

